



Operating Systems CS F372

Threads

BIJU K RAVEENDRAN

Threading Issues

- The fork and exec system calls
 - If one thread in a system calls `fork()`
 - The new process duplicates all threads
 - The new process duplicates only the calling thread.
- Cancellation
 - Task of terminating the thread before it has completed.
 - Cancellation of target thread (the thread that is to be cancelled) may occur in 2 different scenarios
 - Asynchronous cancellation
 - terminates the target thread immediately
 - Deferred cancellation
 - allows the target thread to periodically check if it should be cancelled





```
#include<pthread.h>
```

```
#include<stdio.h>
```

```
#include<asm/unistd.h>
```

```
void *runner(void *param);
```

```
int main(int argc,char *argv[])
```

```
{ pthread_t tid,tid1;
```

```
pthread_attr_t attr;
```

```
pthread_attr_init(&attr);
```

```
pthread_create(&tid,&attr,runner,argv[1]);
```

```
pthread_create(&tid1,&attr,runner,argv[2]);
```

```
printf("1st thread ID=%u & 2nd thread ID=%u\n",tid,tid1);
```

```
if(!fork())
```

```
{ printf("Child PID=%d, PPID=%d\n",getpid(),getppid());
```

```
printf("Child TID=%d, PID=%d\n",syscall(__NR_gettid),getpid());
```

```
}
```

```
#include<unistd.h>
```

```
#include<sys/types.h>
```




else

{

```
wait(NULL);
```

```
printf("Parent:PID=%d,PPID=%d\n",getpid(),getppid());
```

```
printf("Parent:TID=%d, PID=%d\n",  
      syscall(__NR_gettid),getpid());
```

```
}
```

```
pthread_join(tid,NULL);
```

```
pthread_join(tid1,NULL);
```

```
return 0;
```

```
}
```



```
// runner function
```

```
void *runner ( void *param )
```

```
{  
    int upper=atoi(param);  
    int i;  
    int sum=0;  
    if (upper>0)  
    {  
        for ( i=1; i <= upper; i++ )  
        {  
            sum = sum + i; }  
    }
```


```
    printf("From thread:Thread ID=%u,SUM=%d\t PID=%d,  
    PPID=%d\n",pthread_self(),sum,getpid(),getppid());
```

```
    printf("From thread:TID=%d,PID=%d\n",  
    syscall(__NR_gettid),getpid());
```

```
    pthread_exit(0);
```

```
}
```

Windows Multi-threaded C Program



```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 1; i <= Upper; i++)
        Sum += i;
    return 0;
}
```



```
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle,INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n",Sum);
}
```


Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS
- Java threads may be created by:
 - Extending Thread class
 - Implementing the Runnable interface



Linux Threads

- Linux refers to them as *tasks* rather than *threads*
- Thread creation is done through **clone()** system call
- **clone()** allows a child task to share the address space of the parent task (process)
- Flags control behavior

| flag | meaning |
|---------------|------------------------------------|
| CLONE_FS | File-system information is shared. |
| CLONE_VM | The same memory space is shared. |
| CLONE_SIGHAND | Signal handlers are shared. |
| CLONE_FILES | The set of open files is shared. |

Processes & Threads

- Linux uses the same internal representation for processes and threads; a thread is simply a new process that happens to share the same address space as its parent
- A distinction is only made when a new thread is created by the `clone` system call
 - `fork` creates a new process with its own entirely new process context
 - `clone` creates a new process with its own identity, but that is allowed to share the data structures of its parent
- Using `clone` gives an application fine-grained control over exactly what is shared between two threads



Implicit Threading

- Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- Creation and management of threads done by compilers and run-time libraries rather than programmers
- Five methods explored
 - Thread Pools
 - Fork-Join
 - OpenMP
 - Grand Central Dispatch
 - Intel Threading Building Blocks



Signal Handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred
- A signal handler is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled
 - Signal may be handled by
 - A default signal handler
 - A user defined signal handler

Every signal has **default handler** that kernel runs when handling signal - **User-defined signal handler** can override default



Signal Handling

- Options:
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process



Thread Pools

- Create a number of threads in a pool where they await work
- Unlimited threads could exhaust system resources and one solution for this is pooling.
 - Create a number of threads at process startup and place them into a pool. When the server require the thread it is allotted and after completion of task it will be back on pool.
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task

Thread Specific Data

- Allows each thread to have its own copy of data like local variables
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)



Thread Cancellation

- Terminating a thread before it has finished
- Thread to be canceled is **target thread**
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled



Thread Cancellation

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

| Mode | State | Type |
|--------------|----------|--------------|
| Off | Disabled | – |
| Deferred | Enabled | Deferred |
| Asynchronous | Enabled | Asynchronous |

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
 - Cancellation only occurs when thread reaches cancellation point
 - Then cleanup is invoked
- On Linux systems, thread cancellation is handled through signals

Windows Threads

- Windows API – primary API for Windows applications
- Implements the one-to-one mapping, kernel-level
- Each thread contains
 - A thread id
 - Register set representing state of processor
 - Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- The register set, stacks, and private storage area are known as the **context** of the thread



Windows Threads

- The primary data structures of a thread include:
 - ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space
 - KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
 - TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space



Windows Thread Data Structures

